


Shear transformation - the bottom is fixed but the top can move/distract/displace
shear happens in sensor of camera

Transformations in homogenous coordinates



$$\begin{cases} y' = y \\ x' = x + s_{xy} y \end{cases}$$

if the object is lengthier (greater y), then the shear alters the x in proportion to y
y doesn't change.

2D Shear along x relative to the origin:

$$SH_x(s_x) = \begin{bmatrix} 1 & s_x & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} SH & 0 \\ 0 & 1 \end{bmatrix}$$

Transformations in homogenous coordinates

Transformation properties:

1. Combine by matrix multiplication.

2. Preserve the homogeneous coordinate.

3. Inverted by matrix inversion.

Generally, rotation is done w.r.t origin point

Examples:

• rotation about arbitrary point:

$$R_{p, \theta} = T(p) R_{\theta} T(-p)$$

• scale about arbitrary point:

$$S_p(s_x, s_y, s_z) = T(p) S(s_x, s_y, s_z) T(-p)$$

order of multiplication matters based on what happened first.
RTp means the point p was first translated and then rotated.

Multiply Tp first and then R (Tp)

If we want to rotate by an arbitrary point p, then we move the point p to origin, meaning translate by -p and then rotate and then translate back to the place where the point was

Transformations in homogeneous coordinates

$$(TR)^{-1} = R^{-1} T^{-1}$$

this is the formula for inverse. we can also interpret it like, if we want to inverse(cancel) the transformation we did, we cancel what we did last first and then the second thing last, means in reverse we have to cancel.

So, if we rotated and translated initially, we write it as TR. When we cancel, we cancel translation first and then rotation ...so it R⁻¹ T⁻¹

$$T^{-1}(t_x, t_y, t_z) = T(-t_x, -t_y, -t_z)$$

$$S^{-1}(s_x, s_y, s_z) = S(1/s_x, 1/s_y, 1/s_z)$$

$$R_u^{-1}(\theta) = R_u(-\theta) = R_u^T(\theta)$$

For, rotation inverse and transpose are same, because it is an orthogonal matrix

Canceling means undoing the action.



$$R^T R = I$$

Transformations between coordinate systems

Let (x_w, y_w, z_w) be a world coordinate system. Let (x_c, y_c, z_c) be a camera coordinate system. Assume that the camera coordinate system is translated by t and rotated by R with respect to the world coordinate system.


A point $p^{(w)}$ in world coordinates is given by $p^{(c)}$ in camera coordinates. $p^{(c)}$ is related to $p^{(w)}$ by:

$$p^{(c)} = M_{c \leftarrow w} p^{(w)} \quad \text{The world and camera coordinate system are different. if we want to change from world to camera coordinates, then we use this formula} \quad (26)$$

$$M_{c \leftarrow w} = (T(t)R)^{-1} \quad (27)$$

$$= R^{-1}T^{-1}(t) \quad (28)$$

$$= R^T T(-t) \quad (29)$$



$M_{c \leftarrow w}$: align camera with world (cancel translation then cancel rotation)

How did the camera coordinate get there???--->it was at the origin of the world, it was then rotated and translated w.r.t the world coordinate system. Now, we wanted to align the world origin/points with whom as base the camera was rotated to the camera's current origin. So, it actually means to be reversing what we did before.

If we want to reverse it, then cancel translation first and then rotation to get back the world origin and camera origin... same as the point where the previous world coordinate was x_c, y_c, z_c --->So, it looks like as though we are moving the camera back to the world origin 0,0,0 at the end because whatever(world) was used as base to move the camera from the origin of the world to some point of the world, Now we are actually moving the world itself to get back to the origin of the camera.

Transformations between coordinate systems

The rotation R is given by:

taking the unit vectors of the camera coordinate s/y gives the Rotation matrix that was rotation initially done by camera from the origin of the world coordinate system to the camera coordinate system

$$R = \begin{bmatrix} \hat{x}_c & \hat{y}_c & \hat{z}_c \end{bmatrix}$$


The inverse transformation is given by:

$$M_{w \leftarrow c} = (M_{c \leftarrow w})^{-1}$$

$$= (R^T T(-t))^{-1}$$

$$= T(t)R$$

General camera model



General camera model

P is 3dH and p is 2dH

Perspective projection in camera coordinates:

$$p^{(c)} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} P^{(c)} = K[I|0]P^{(c)} \quad (34)$$

The perspective projection we want:

$$\text{We want } p^{(i)} = QP^{(w)} \quad (35)$$

$$p^{(i)} = M_{i \leftarrow c} p^{(c)} \quad (36)$$

$$= M_{i \leftarrow c} K[I|0]P^{(c)} \quad (37)$$

$$= M_{i \leftarrow c} K[I|0]M_{c \leftarrow w} P^{(w)} \quad (38)$$

Assuming that the camera is translated by t and rotated by R with respect to the world ($R = \begin{bmatrix} \hat{x}_c & \hat{y}_c & \hat{z}_c \end{bmatrix}$):

perspective projection--a technique used to project three-dimensional points onto a two-dimensional plane, mimicking how the human eye perceives depth and distance.

$p(c) = (X', Y', Z')$ represent the transformed position of a 3D point in the camera's perspective, where X' and Y' are used for mapping to the 2D image plane, and Z' provides depth information essential for understanding the spatial relationship of objects in the scene.

when we divide by Z' , we get the 2d point, However we should take into account, the values of $k_u, k_v, u_0, v_0 \rightarrow$ and that is why we have the extra matrix $M_{i \leftarrow c}$

General camera model

As detailed in before slides

$$\begin{aligned} M_{c \leftarrow w} &= R^T T(-t) \\ &= \begin{bmatrix} R & 0 \\ 0 & 1 \end{bmatrix}^T \begin{bmatrix} I & -t \\ 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} R^T & -R^T t \\ 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} R^* & T^* \\ 0 & 1 \end{bmatrix} \end{aligned}$$

The rotation/translation of the world with respect to the camera are extrinsic camera parameters:

Camera calibration- generally returns R^* and T^*

$$\begin{aligned} R^* &= R^T \\ T^* &= -R^T t \end{aligned}$$

While R and T are intrinsic camera parameters where the rotation/translation of the camera happens w.r.t to origin of the world

General camera model

Intrinsic camera parameters:



notation	meaning	units
k_u	scale in x relating pixels to mm	pixels/mm
k_v	scale in y relating pixels to mm	pixels/mm
u_0	translation of the principal point in x	pixels
v_0	translation of the principal point in y	pixels
f	focal length	mm
k_u, k_v	focal length	pixels

k_u = number of pixels/mm in x
 k_v = number of pixels/mm in y
 if the pixels are square, then $k_u = k_v$

The transformation $M_{i \leftarrow c}$ is composed of scale and translation:

$$M_{i \leftarrow c} = \begin{bmatrix} k_u & 0 & u_0 \\ 0 & k_v & v_0 \\ 0 & 0 & 1 \end{bmatrix}$$

u_0 and v_0 corresponds to the translation of the center of the image w.r.t to the camera coordinate system

u_0 and v_0 are the same for an ideal camera.

For a 1000x1000 image, $u_0=v_0=500$.

However, if there is a defect in the camera, it would be different.



While going from camera to image coordinates, there is no rotation, but there is translation, because the origin of the camera coordinate system and image center point vary. There is also scaling transformation as camera coordinate is in mm and image is in pixel.

General camera model

$$K = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Combining the transformation matrices:

$$\begin{aligned} p^{(i)} &= M_{i \leftarrow c} K [f|0] M_{c \leftarrow w} P^{(w)} \\ &= K^* [f|0] \begin{bmatrix} R^* & t^* \\ 0 & 1 \end{bmatrix} P^{(w)} \\ &= K^* [R^* | t^*] P^{(w)} \end{aligned}$$

This is the perspective projection that also includes the transformations in the coordinate s/y

The matrix K^* contains the intrinsic camera parameters:

$$\begin{aligned} K^* &= M_{i \leftarrow c} K \\ &= \begin{bmatrix} k_u & 0 & u_0 \\ 0 & k_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} f k_u & 0 & u_0 \\ 0 & f k_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} \alpha_u & 0 & u_0 \\ 0 & \alpha_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \end{aligned}$$

The parameters α_u, α_v specify scale in pixels.

Look at example sum in word document

General camera model

Shear is the smallest distortion that happens from the sensors. It is usually acceptable but for important tasks that need high precision and care, it needs to be corrected.

To correct this, we add the additional terms in the K^* term

When allowing for shear:

$$\begin{aligned} K^* &= \begin{bmatrix} \alpha_u & \alpha_u \tan \alpha & u_0 \\ 0 & \alpha_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} \alpha_u & s & u_0 \\ 0 & \alpha_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \end{aligned}$$

s is almost 0, it is very small value

The parameter s is skew (in pixels).

There may not be this big a diff but there may be 0.1 pixel diff - but accuracy matters in some sensitive tasks



Radial lens distortion

Refer diagram from notes doc

$$p^{(i)} = \begin{bmatrix} \lambda & 0 \\ 0 & \lambda \\ 0 & 0 & 1 \end{bmatrix} K^* [R^* | T^*] P^{(w)}$$

$$\lambda = 1 + k_1 d + k_2 d^2$$

linear distortion coefficient
quadratic distortion coefficient


d = distance from center

larger shrink away from the center

k_1 and k_2 depends on the camera calibration

wide angle lens usually creates distortion
Eg - security cameras, classroom cameras

More distortion seen away from the center of the image



The lines are supposed to be parallel as per real world but distorted due to wide angle lens.

the lambda value depends on the distance (d) between the distorted point and the center of the image.

Here, we again have division and hence it is non linear. So we cannot use it directly. Instead, we transform the distorted image to the original (undistorted world type) image using warp, then we use that image for further computation

there is a warp function to do this

Radial lens distortion


Warp the image to correct distortion
(warp using estimated distortion parameters)



Weak perspective camera


Perspective

weak perspective



Weak perspective camera

Weak perspective is correct when depth variation in the scene is small compared with distance from camera



$$e = |M_w \mathbf{p} - M \mathbf{p}| = \frac{\Delta}{d_0} (\underbrace{MP - P_0}_{\substack{\text{depth variation} \\ \text{distance from center} \\ \text{distance from camera}}})$$

placing the object far from the camera, makes it a weak perspective camera

Affine is not real but used for simplification

Affine camera

$$M_{\text{affine}} = \begin{bmatrix} a & b & c & d \\ e & f & g & h \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

computational model

used to convert 3dH(4x1) to 2dH (3x1)